

Measuring and Profiling — Notes and Examples

Slide 6: Profiling - `utils::Rprof()`

- The function call stack is a record of the function currently executing, the function that called the function, and so on (from right to left)
- The default time interval between sampling is 0.02 seconds. If the profiled code executes faster you need to pass an appropriate sampling time to the argument `interval`.
- The output of `Rprof()` is a complete listing of the function call stack at every sampling iteration (which, on its own, is often not very informative) which is printed to a binary file in the current working directory

Example: profiling a call of `replicate()` — `ctd`.

- A more useful output is produced by `utils::summaryRprof()`:

By appending `$by.total` we get the time spend in each function by the total run time.

```
summaryRprof(tmp)$by.total
```

So `rnorm()` is six levels deep in the call stack and it seems that R spends most of time evaluating `rnorm(1e6)` (which takes ~ 0.3 seconds).

- By appending `$by.self` the results are adjusted for the time needed to run functions above the current function in the call stack.

```
summaryRprof(tmp)$by.self
```

As expected, random number generation takes ~ 100% of the total computation time.

Slide 7: Visualising profiles: `profvis::profvis()`

- We use `profvis::pause()` since the time spend in `sys.sleep()` is not measured as computing time.
- Visualisation using `profvis()` works best when the code is sourced from an `.R`-script:

```
# source f(), g(), h() form R script
source("codes/profiling-example.R")

# visualise profiling results
profvis(f())
```

- After profiling, `profvis()` opens an interactive HTML window will *RStudio* which lets us explore the results (see Figure 1).
- The interface provided by `profvis()` connects the profiling data back to the source code. This makes it easier to build up a ‘mental’ model of what you may improve.

The top pane in the left panel of Figure 1

- displays bar plots of running times and memory allocations with the latter being no issue here (why?)
- provides a good overall feel for bottlenecks but is quite imprecise about the *cause*

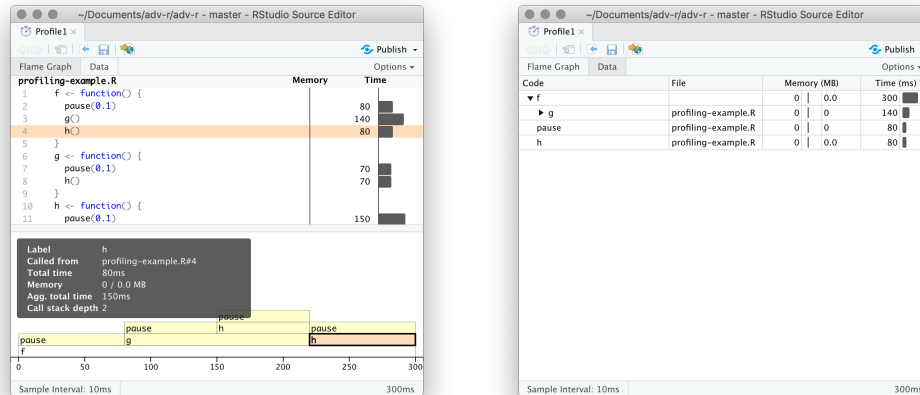


Figure 1: `profvis()`

`h()` is not significantly slower than `g()`: `h()` is reported to take 150 ms (twice as long as `g()`) because it is called two times.

The bottom pane shows a *flame graph* of the full call stack. The full sequence confirms that

- `h()` is called from two different places (once by `g()` and once by `f()`)
- the execution time of `h()` is roughly the same in each call.
- Mousing over a cell in the function stack indicates the corresponding line in the source code.

The data tab (Figure 1, right panel) provides a tree-based representation of the top pane. This is useful for analysing more complicated components of the code.

Slides 8: Visualising profiles: `profvis::profvis()`

- Notice that the actual computation of model is relatively fast. The plot functions are the root of all evil!

Slide 11: Memory Profiling

Solution to Exercise

What is going on?

- It looks like R spends most of the time modifying the data in-place, but that's not actually what's happening internally.
- The memory column indicates that large amounts of memory are being allocated (right bar) and freed (left bar)

Reasons:

1. A new memory object is generated by modifying a copy of the 'old' `x` which is then reassigned to `x` in each iteration
 2. Garbage collection (GC) automatically frees memory by deleting no more required objects
- While `c()` runs for a total of 170 ms, a considerable amount of this time is due to garbage collection (`<GC>`)

When you see the garbage collector taking up a lot of time, you can often come up with a more efficient alternative.

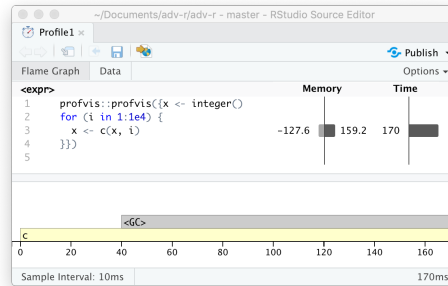


Figure 2: Memory profiling with profvis()

Slide 12: Memory Profiling

Solution to Exercise

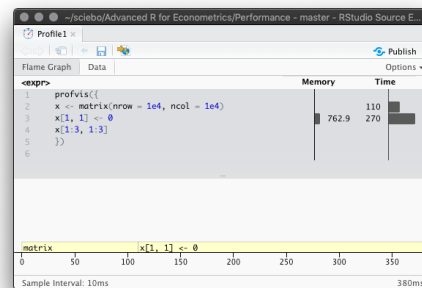


Figure 3: profiling coercion to another type

What is going on?

The code seems fairly innocent, but it turns out that it is very inefficient in terms of both memory and speed:

- Line 3 assigns a numeric value upon which **x** is written to the memory. It requires $(1e4 * 1e4 * 8)/1024^2 = 762.9395$ MB.
- **x** is initialized with type **logical** (it is filled with NAs) and thus cannot contain numeric values.
- **x[1, 1] <- 0** internally coerces **x** to a **numeric** matrix before assigning 0 to the (1,1) element which is quite costly: initializing **x** was more than twice as fast!

Especially if a large number of computations need to be performed, coercion should be avoided wherever possible!

Slide 14: Memory Profiling

Example: primitive functions

Primitive functions cannot be profiled:

```
profvis({
  sqrt(sum(abs(rnorm(sum(1e6))))))
})
```

Slide 21:

```
set.seed(42)

create_df <- function(rows, cols) {
  as.data.frame(matrix(
    unlist(replicate(cols, runif(rows, 1, 1000), simplify = FALSE)),
    nrow = rows, ncol = cols))
}

results <- bench::press(
  rows = c(10000, 100000),
  cols = c(10, 100),
  {
    dat <- create_df(rows, cols)
    bench::mark(
      min_iterations = 100,
      bracket = dat[dat$x > 500, ],
      which = dat[which(dat$x > 500), ],
      subset = subset(dat, x > 500)
    )
  }
)
```

ggplot2::autoplot() automatically generates a facet plot for results.

```
plot(results)
```

Slide 22: Microbenchmarking — Exercises

1. Instead of using `bench::mark()`, you could use the built-in function `system.time()` which is, however, much less precise, so you'll need to repeat each operation many times with a loop, and then divide to find the average time of each operation, as in the code below.

How do the estimates from `system.time()` compare to those from `bench::mark()`? Why are they different?

Solution:

As `bench::mark()` doesn't calculate the mean value, we calculate it from the time list-column in the tibble output.

```
n <- 1e6
x <- runif(100)

bench_res <- bench::mark(
  sqrt(x),
  x ^ 0.5
)

# Compute mean across all runs
t_sqrt_bench <- mean(unlist(bench_res[1, "time"]))
t_power_bench <- mean(unlist(bench_res[2, "time"]))

t_sqrt_systemtime <- system.time( for (i in 1:n) sqrt(x) ) / n
t_power_systemtime <- system.time( for (i in 1:n) x^0.5 ) / n
```

```

# Compare the results
t_sqrt_systime["elapsed"]
#> elapsed
#> 1.06e-06
t_sqrt_bench
#> [1] 8.97e-07

t_power_systime["elapsed"]
#> elapsed
#> 8.64e-06
t_power_bench
#> [1] 1.05e-05

```

Both approaches get the order of magnitude right. The results differ a little and `bench::mark()` is generally more precise.

- Here are two other ways to compute the square root of a vector. Which do you think will be fastest? Use microbenchmarking to test your answers.

Solution:

```

x <- runif(100)

bench::mark(x^(1/2),
            exp(log(x)/2),
            relative = TRUE)

```

So $x^{1/2}$ is faster (which is due to overhead produced by calling `exp()` and `log()`).

Slide 25: (Why) is R slow? — Slow Code Interpretation

In the example, the interpreter cannot predict that the loop always adds 1 to an integer `x`. R needs to look for the right `+` method (the method for adding two doubles) in *every* iteration of the loop!